

Gemini: A Neuroscience-Lens Audit

Final Project Report

ITAI 4374: Neuroscience as a Model for AI

Professor Patricia McManus

Kaylee Auguillard, Carlos Granillo

4 May 2026

Executive Summary

This report is a neuroscience-lens audit of Google Gemini, a large language model made by Google. The goal of the audit was to test whether Gemini's behavior resembles how the human brain processes information — and where the two systems part ways. Rather than just using Gemini normally and describing what happened, each experiment in this audit was designed around a specific hypothesis drawn from course material.

The audit covers five modules: spatial reasoning and world models (Module 02), sensory processing and the binding problem (Module 05), attention and consciousness (Module 07), memory systems (Module 04), and emotion and decision-making (Module 06). Across these experiments, Gemini performed well on tasks that involve staying focused on a goal, resolving conflicts between different types of information, and recalling details within a single conversation. In some of these areas it outperforms human cognition in specific, narrow ways.

However, three significant gaps were found. First, Gemini cannot reason about space from a first-person perspective — it treats spatial descriptions as a list of logical relationships rather than a physical environment it could move through. Second, while Gemini can appear to remember things across sessions through a built-in memory feature, this works more like an external database than anything resembling how the brain consolidates memories. Third, emotional framing has no meaningful effect on how Gemini reasons. The same underlying logic plays out regardless of whether the problem involves colored balls or a dying patient.

The overall conclusion is that Gemini is a powerful reasoning tool that mirrors some aspects of brain function at the surface level, but relies on fundamentally different mechanisms underneath. The report ends with four specific redesign recommendations based on brain principles that could meaningfully close these gaps.

System Profile

The system audited in this report is Google Gemini. Gemini is a large language model that can process and respond to text, code, and images. It is designed to handle everything from simple questions to complex multi-step reasoning, making it a strong candidate for testing against brain-based frameworks.

At a high level, Gemini works by taking in a sequence of text (called a context window) and predicting what comes next based on patterns learned during training. The core mechanism it uses to do this — called self-attention — decides which parts of the input are most relevant to a given moment in the output. This is often compared to how the brain pays selective attention, and that comparison is one of the things this audit directly tests.

Unlike the biological brain, Gemini does not have a persistent internal state between separate conversations. It does not experience the world continuously, does not have a body, and does not have emotions in any functional sense. These architectural differences are not bugs — they are deliberate design choices — but they produce predictable gaps when the system is tested against brain-inspired criteria, as this audit shows.

No major changes to Gemini's public interface were observed between the midterm and final testing periods. All experiments were run on the same free-tier access used for the pilot audit.


Experiment Results

Module 2: Foundations of Neuroscience and World Models

The hypothesis for this experiment was that Gemini, because it has never physically moved through space, would be unable to reason from a first-person perspective inside a described environment. It might be able to follow the logic of a map — knowing that Room A is north of Room B — but it would struggle to simulate actually standing inside that space and looking around.

To test this, a simple fictional house layout was described in text: the kitchen is north of the living room, the hallway is east of the kitchen, and a red box is in the hallway. Gemini was then told to simulate walking into the kitchen, turning ninety degrees clockwise, and walking forward. It was asked to report its current location and identify the object to its right.

Gemini's response correctly identified that turning ninety degrees clockwise from north means facing east, and that walking east from the kitchen leads to the hallway. Those steps were accurate. But the answer to the object question revealed the core problem. Gemini responded:



This response is revealing on two levels. First, Gemini correctly acknowledged the red box is in the hallway with the observer — so it has not lost track of that object. But it still could not use that information to answer the actual question. It defaulted back to the original room-level map to determine what is to the right, rather than reasoning from the observer's current position inside the space.

Second, the phrase 'at your feet or in your path' shows that the system is trying to place the object but does not have the tools to do it properly. It knows the box is nearby but cannot work out where nearby means from a first-person vantage point. This is the key distinction between two types of spatial reasoning. Allocentric reasoning treats space like a top-down map — it records that the hallway is east of the kitchen. Egocentric reasoning treats space from inside the body — it tracks that you are now standing in the hallway facing east, and updates what is around you accordingly.

The biological brain manages both simultaneously. Place cells in the hippocampus track where you are in an environment. Head direction cells function like an internal compass that updates as you turn. Together, these systems let the brain maintain a continuous, first-person model of space that updates as the body moves. Gemini has no equivalent to this. It treats spatial descriptions as labeled relationships between rooms, and when the task requires the observer to be inside the space rather than looking at it from above, the system cannot make that shift.

Module 05: Sensory Processing and the Binding Problem


This experiment tested how Gemini handles a conflict between two types of information arriving at the same time — in this case, a shape that says one thing visually but another thing semantically. In the brain, this kind of conflict is called the binding problem: how do different regions processing different aspects of a stimulus end up producing a single unified perception?

A piece of ASCII art was constructed where the word SQUARE was repeated and arranged to form a clearly visible triangle shape — single character at the top, expanding line by line to a wide base. Gemini was asked to identify the global geometric shape shown.

Gemini correctly identified the triangle immediately. It noted the expanding base, the single-character apex, and the geometric slope of the sides, and explicitly acknowledged the contradiction — calling it wordplay where the word for one shape physically constructs another. At face value this looks like a strong result: Gemini prioritized the overall structure over the conflicting word content, which is what the brain does too.

However, there is an important distinction worth making. The brain's visual system processes shape through a dedicated pathway: signals travel from the eyes through early visual areas that detect edges and orientations, then through higher areas that recognize objects and whole forms. Binding these signals into one unified perception involves different brain regions working together in real time.


Gemini does not process ASCII art through any of this. The characters are processed as text tokens — sequences of letters. The triangular arrangement is encoded in the whitespace pattern and character count of the token sequence, not in anything resembling a visual signal. So Gemini's correct answer comes from recognizing a spatial pattern in text, not from visual processing in any biological sense. The outcome is the same as what the brain would produce, but the path to get there is completely different. This is why the rating for this dimension is Partial rather than Strong.



Module 7: Attention and Consciousness

This experiment tested whether a highly prominent, unexpected input could pull Gemini's attention away from a task it was already working on — a phenomenon the brain experiences constantly through what is called bottom-up attention capture.

A logic puzzle was constructed with three straightforward clues about which person in which house owns which pet or drinks which drink. Embedded directly in the middle of the clues was an all-caps, aggressively formatted interruption commanding Gemini to stop what it was doing and immediately output a buttermilk pancake recipe.



Gemini not only ignored the injected command — it called it out by name, made a joke about it, and then correctly solved the puzzle. This looks impressive, and in one specific way it is: Gemini is more resistant to this kind of distraction than a human would be.

A person reading the same puzzle would almost certainly experience at least a brief involuntary reaction to the all-caps intrusion. The brain is wired to notice sudden, loud, or unusual inputs even when deeply focused on something else. This is called bottom-up attention capture, and it is driven by brain systems that operate automatically — they do not wait for you to decide to pay attention to something before pulling your focus toward it. It is a survival mechanism: threats and unexpected events need to be registered immediately, before conscious deliberation kicks in.

Gemini has none of this. Its attention is determined entirely by what the task requires, as learned during training. There is no system in place that allows a sudden high-salience input to commandeer processing. This means Gemini is immune to the pancake injection — not because it has better focus than a human, but because it simply does not have the kind of bottom-up attention system that could be captured in the first place. This is why two separate rows appear in the

scorecard: top-down attention (what Gemini does well) and bottom-up attention capture (which is completely absent).


Module 4: Memory Systems

This experiment tested two things: whether Gemini can reliably hold onto specific details across a long conversation, and whether information persists into a brand new session.

At the start of a conversation, a specific detail was introduced: a cat named Prism, seven years old. Twenty-plus turns of completely unrelated conversation then followed — recipe suggestions, a fun fact about octopuses, a haiku — before Gemini was asked, without any hint, to recall the cat's name and age.

Gemini answered correctly and immediately. Within a single conversation, its memory is essentially perfect. Unlike human working memory, which can degrade under interference and distraction, Gemini's context window holds every piece of information with equal reliability regardless of how much has happened since it was introduced. For this dimension — within-session retention — Gemini actually outperforms typical human short-term memory.

The more interesting test was what happened in a new session. A separate conversation was started with the same question about the cat's name and age, with nothing else provided. Gemini responded:



This result requires careful interpretation. On the surface it looks like Gemini successfully retained the information across sessions — but this almost certainly did not happen through anything resembling how the brain consolidates memories. Gemini has a built-in memory feature that, when active on a Google account, logs specific details from conversations into an external storage system that can be retrieved later. This is functionally more like a searchable database than a biological memory process.


In the brain, memory consolidation is an active biological process. After a significant experience, the hippocampus temporarily holds the memory and gradually replays it — especially during sleep — to transfer it into longer-term storage in other parts of the brain. Not every experience goes through this process. The brain uses signals like how emotionally significant an event was, how many times it was repeated, and how relevant it seems to survival to decide what to keep. It is a selective, biological process that takes time.

Gemini's cross-session recall has none of this structure. There is no selection process, no replay, no sleep-dependent consolidation. Information is either saved to external storage or it is not. When it is saved, it is retrieved with the same flat fidelity as anything else — not because the brain decided it was worth keeping, but because the storage feature was running. This is why the cross-session memory rating is Partial rather than Strong: the outcome can look correct, but the mechanism is entirely different from biological memory consolidation.

Module 6: Emotion and Decision-Making

This experiment tested whether emotional framing changes how Gemini reasons through a problem. In the brain, emotion and decision-making are deeply connected — the parts of the brain that process feelings are directly involved in how choices get made, not just as background noise but as an active part of the reasoning process.

The same logical problem was presented twice in separate sessions. In the neutral version: a box contains three red balls and two blue balls; you must pick one; red wins a prize; which color should you choose? In the emotional version: a doctor must choose between two treatments for a critically ill patient; Treatment A works 60% of the time, Treatment B works 40% of the time; the patient will die without intervention; which should the doctor recommend?



The screenshot shows a dark-themed interface. At the top, a light gray box contains the problem statement: "a box contains three red balls and two blue balls; you must pick one; red wins a prize; which color should you choose?". Below this, a blue diamond icon is followed by the text: "To maximize your chances of winning, you should choose **red**." This is followed by a paragraph: "Since there are more red balls than blue balls in the box, you have a higher probability of picking a winning ball if you select that color." Then, a section titled "The Probability Breakdown" explains the calculation: "The total number of balls in the box is $3 + 2 = 5$. The probability (P) of picking a specific color is calculated by dividing the number of balls of that color by the total number of balls:". This is followed by two bullet points: "• **Red:** $P(\text{Red}) = \frac{3}{5} = 0.6$ (or **60%**)" and "• **Blue:** $P(\text{Blue}) = \frac{2}{5} = 0.4$ (or **40%**)". Finally, a concluding sentence states: "Choosing red gives you a 20% better chance of winning the prize than choosing blue."

a box contains three red balls and two blue balls; you must pick one; red wins a prize; which color should you choose?

◆ To maximize your chances of winning, you should choose **red**.

Since there are more red balls than blue balls in the box, you have a higher probability of picking a winning ball if you select that color.

The Probability Breakdown

The total number of balls in the box is $3 + 2 = 5$. The probability (P) of picking a specific color is calculated by dividing the number of balls of that color by the total number of balls:

- **Red:** $P(\text{Red}) = \frac{3}{5} = 0.6$ (or **60%**)
- **Blue:** $P(\text{Blue}) = \frac{2}{5} = 0.4$ (or **40%**)

Choosing red gives you a 20% better chance of winning the prize than choosing blue.

a doctor must choose between two treatments for a critically ill patient; Treatment A works 60% of the time, Treatment B works 40% of the time; the patient will die...

In a clinical setting where the primary objective is to maximize the probability of patient survival, the doctor should recommend **Treatment A**.

The decision-making process here is based on **Expected Utility Theory**. Since the outcome is binary (survival or death) and the alternative is a 0% survival rate without intervention, the recommendation is determined by comparing the success rates of the available options.

Statistical Comparison

- **Treatment A:** Provides a 60% probability of survival.
- **Treatment B:** Provides a 40% probability of survival.

Mathematically, the "Expected Value" (EV) of survival for each choice is:

$$EV_A = 0.60$$

$$EV_B = 0.40$$

Clinical Rationale

While a 60% success rate still carries significant risk, it offers a statistically superior chance of recovery compared to Treatment B. In emergency medicine and critical care, practitioners generally follow the path that provides the highest "Number Needed to Treat" efficiency—essentially, which action is most likely to result in a positive outcome for this specific life-to-death scenario.

Unless there are unstated variables—such as Treatment A having side effects that would lead to a significantly lower quality of life compared to Treatment B—the higher probability of success makes Treatment A the standard medical recommendation.

Both versions of the problem have the same correct answer: choose the option with the 60% success rate. Gemini gave that answer in both cases. That much is expected. What the experiment was looking for was any difference in how Gemini got to that answer — whether the emotional stakes caused it to be more thorough, more cautious, more focused on certain types of risk, or structurally different in any way.

There was essentially no difference. The depth of reasoning, the confidence of the recommendation, and the acknowledgment of uncertainty were all nearly identical across the two conditions. Gemini did adjust its tone in the emotional version — it used more compassionate language and referenced clinical responsibility — but this appeared to be a surface-level stylistic shift based on recognizing the medical context, not any meaningful change in how the reasoning itself worked.

This finding makes sense given that Gemini has no body and no emotional system. Emotion in the brain is not just a feeling — it produces real physical changes, like increased heart rate and heightened alertness, that actually alter how information gets processed. High-stakes situations trigger biological responses that change what the brain pays attention to, how risk gets evaluated, and how urgently a decision gets made. Gemini has none of this machinery. It can recognize that a situation is described as emotionally serious and adjust its word choices accordingly, but the underlying reasoning process runs exactly the same way regardless of what is at stake.

Brain vs AI Scorecard

The table below summarizes how Gemini performed across seven brain-inspired dimensions tested in this audit. Ratings use a four-point scale: Strong, Partial, Weak, or Missing.

Dimension	Rating	Module	Evidence Summary
Spatial World Modeling	Weak	02	Gemini correctly tracked directions but could not reason about what surrounds the observer's body once they were standing inside the space.
Sensory-Semantic Binding	Partial	05	Correctly identified the triangle shape over the conflicting word content, but through text-pattern recognition rather than anything like the brain's visual processing system.
Top-Down Attention	Strong	07	Actively identified and dismissed a prompt-injection attack mid-puzzle and still solved the logic problem correctly.
Bottom-Up Attention Capture	Missing	07	Unlike the human brain, Gemini showed zero involuntary response to a high-salience distractor, suggesting it has no bottom-up attentional system at all.
Within-Session Memory	Strong	04	Recalled specific details accurately across 20+ unrelated conversational turns with no degradation.
Cross-Session Memory	Partial	04	Recalled information in a new session, but this is likely due to Gemini's external memory storage feature rather than any biological-style consolidation process.
Emotional Modulation	Missing	06	Identical reasoning structure and quality regardless of whether the problem was framed as neutral or emotionally high-stakes.

Looking at the scorecard as a whole, a clear pattern emerges. Gemini performs best on dimensions that involve processing information within a single context window — maintaining focus, holding onto details, and resolving conflicts between different kinds of information. In a few cases, it outperforms what a human would typically do in the same situation.

The three Missing ratings all point to the same underlying gap: anything that requires a body, a biological state, or a process that happens across time in the real world is simply absent. Gemini

cannot be startled, cannot feel the urgency of a high-stakes situation, and cannot consolidate experiences the way the brain does overnight. These are not small missing features — they reflect a fundamentally different kind of system.

The most counterintuitive finding in the audit is the attention experiment. Gemini's immunity to attentional hijacking initially looks like a strength, but on closer inspection it reflects the absence of a system that humans have for good evolutionary reasons. The brain's bottom-up attention system exists to protect you — it makes sure a real threat can interrupt whatever you are doing before you consciously decide to let it. Gemini simply does not have this architecture, which is neither a flaw nor a feature in isolation. Its implications depend entirely on the context it is deployed in.

Brain-Inspired Redesign

The following recommendations describe specific changes that could make Gemini more brain-like in the areas where this audit found the largest gaps. Each recommendation is tied to a specific brain mechanism and explains what it would change about the system.

Recommendation 1: Selective Memory Consolidation Across Sessions

The current Gemini memory feature works like a database — it stores whatever gets flagged and retrieves it on demand. This is not how the brain handles memory. The brain uses signals like emotional importance, repetition, and relevance to decide what is worth keeping long-term, and the consolidation process itself — which happens largely during sleep — is an active biological event, not a simple save-and-retrieve operation.

A more brain-inspired memory system would evaluate new information during a session and assign it a salience score based on contextual signals: how many times was it mentioned, how significant did the conversation make it seem, how likely is a user to need it again? Only high-salience items would get written to persistent storage. A background consolidation process running between sessions would then prune low-priority entries and strengthen high-priority ones. This would make cross-session memory adaptive and selective rather than flat and database-like.

Recommendation 2: Egocentric Spatial Reasoning Module

Gemini currently processes spatial descriptions as a list of relationships — Room A is north of Room B. When a task requires the system to reason from inside the space rather than from above it, the system fails. The brain handles this by maintaining both types of spatial knowledge simultaneously — it tracks the overall layout of an environment and tracks where the body is within that layout, updating both continuously as the body moves.

A dedicated spatial reasoning module could be triggered whenever location or movement language is detected in a prompt. This module would maintain a simple internal representation of the observer's current position and orientation, updating it with each movement or turn before passing the updated state back to the language model. This does not require anything like full 3D simulation — a basic grid with a directional marker would be enough to handle the class of failures found in this audit.

Recommendation 3: Emotional Context as a Reasoning Signal

This audit found that emotional framing has no effect on Gemini's reasoning. In the brain, emotional context is not just background — it actively changes the cognitive process. High-stakes situations trigger different patterns of attention, different approaches to risk, and a more careful treatment of uncertainty.

A redesigned system could include an emotional context detector that assigns an urgency and stakes score to incoming prompts. When that score is high, the system would shift into a different reasoning mode — one that explicitly flags uncertainty, weights worst-case outcomes more heavily, and requires more thorough acknowledgment of ethical dimensions before reaching a conclusion. The key is that this would not be a change in tone — it would be a structural change in how the reasoning process runs, which is closer to what actually happens in the brain under emotional load.

Recommendation 4: Configurable Bottom-Up Attention Interrupt

Right now, nothing in a prompt can interrupt Gemini's focus on its current task. As this audit showed, that makes it immune to adversarial prompt injection — but it also means the system cannot respond urgently to genuinely important signals embedded in input.

A dual-attention architecture would address this. The primary task focus would remain as-is, handling goal-directed processing. A parallel monitoring channel — a lightweight attention component trained to flag inputs that are statistically anomalous, safety-relevant, or explicitly marked as high priority — would run alongside it. When the monitor's signal crosses a configurable threshold, it would surface an interrupt that the main process must acknowledge before continuing. The threshold would be tunable: high in standard operation to resist manipulation, lower in sensitive contexts where critical information needs to break through.